

Interactions (V)

Last time...

Factorial ANOVA

- Estimated marginal means
- Estimated cell means
- Sums of squares

Power

- Omnibus test
- Coefficients

ANOVA vs Regression

Factorial ANOVA

Interaction tests whether there are differences in differences.

Simple main effect -- the effect of Factor A at a specific level of Factor B

Regression

Interaction tests whether slope changes.

Simple slopes -- the slope of Variable A at a specific level of Variable B

Power

The likelihood of finding an effect *if the effect actually exists.*

When calculating power for the omnibus test, use the expected multiple R^2 value to calculate an effect size: $f^2 = \frac{R^2}{1-R^2}$

To estimate power for a single coefficient, you need to consider (1) how much variance is accounted for by just the variable and (2) how much variance you'll account for in Y overall.

$$f^2 = \frac{R_Y^2 - R_{Y.X}^2}{1 - R_Y^2}$$

Effect sizes (interactions)

To start our discussion on powering interaction terms, we need to first consider the effect size of an interaction.

How big can we reasonably expect an interaction to be?

- Interactions are always partialled effects; that is, we examine the relationship between the product of variables X and Z only after we have controlled for X and controlled for Z. How does this affect the size of the relationship between XZ and Y?

Effect sizes (interactions)

The effect of XZ and Y will get **smaller** as X or Z (or both) is related to the product

The semi-partial correlation is always smaller than or equal to the zero-order correlation.

McClelland and Judd (1993)

Is it more difficult to find interaction effects in experimental studies or observational studies?

What factors make it relatively easier to find interactions in experimental work?

Influencing power in experimental studies

- No measurement error of IV
 - don't have to guess what condition a participant is in
 - measurement error is exacerbated when two variables measured with error are multiplied by each other
- Experimentalists are more likely to find cross-over interactions; observational studies may be restricted to fan interactions
 - cross-over interactions are easier to detect than fan interactions

Influencing power in experimental studies

- Experimentalists can concentrate scores on extreme ends on both X and Z
 - in observational studies, data tends to cluster around the mean
 - increases variability in both X and Z, and in XZ
- Experimentalists can also force orthogonality in X and Z
- Experimentalists can study the full range of X in an experiment

McClelland and Judd's simulation

For the experiment simulations, we used 2×2 factorial designs, with values of X and Z equal to +1 and —1 and an equal number of observations at each of the four combinations of X and Z values.

```
X = rep(c(-1,1), each = 50)
Z = rep(c(-1,1), times = 50)
table(X,Z)
```

```
##      Z
## X    -1   1
##   -1 25 25
##   1  25 25
```

McClelland and Judd's simulation

For the field study simulations, they used values of X and Z that varied between the extreme values of +1 and —1. More specifically, in the field study simulations, values of X and Z were each sampled independently from a normal distribution with a mean of 0 and a standard deviation of 0.5. Values of X and Z were rounded to create equally spaced 9-point scales ranging from -1 to +1 because ranges in field studies are always finite and because ratings are often on scales with discrete intervals.

McClelland and Judd's simulation

For field studies

```
X = rnorm(n = 100, mean = 0, sd = .5)
Z = rnorm(n = 100, mean = 0, sd = .5)
X = round(X/.2)*.2
Z = round(Z/.2)*.2

psych::describe(data.frame(X,Z), fast = T)
```

```
##    vars     n   mean    sd median   min  max range skew kurtosis    se
## X      1 100 -0.05  0.5       0 -1.2 1.2    2.4  0.04    -0.41 0.05
## Z      2 100  0.03  0.5       0 -1.6 1.2    2.8 -0.44      0.63 0.05
```

For both: $\beta_0 = 0$, $\beta_X = \beta_Z = \beta_{XZ} = 1$.
 $N = 100$, and randomly sampled normally distributed errors ($\mu = 0$, $\sigma = 4$)

```
Y = 0 + 1*X + 1*Z + 1*X*Z + rnorm(n = 100, mean = 0, sd = 4)
summary(lm(Y ~ X*Z))
```

```
##
## Call:
## lm(formula = Y ~ X * Z)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -11.2002  -3.2569  -0.4427   3.1085  10.0917
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.7331    0.4710   1.557   0.123
## X            1.3502    0.9430   1.432   0.155
## Z            1.3674    0.9615   1.422   0.158
## X:Z         -2.8471    2.0902  -1.362   0.176
##
## Residual standard error: 4.679 on 96 degrees of freedom
```

```
set.seed(0305)
```

```
# for experimental studies
sim = 100
ebeta_xz = numeric(length = sim)
et_xz = numeric(length = sim)
for(i in 1:sim){
  # simulate data
  X = rep(c(-1,1), each = 50)
  Z = rep(c(-1,1), times = 50)

  Y = 0 + 1*X + 1*Z + 1*X*Z
  rnorm(n = 100, mean = 0,
  #run model
  model = lm(Y ~ X*Z)
  coef = coef(summary(model))
  #extract coefficients
  beta = coef["X:Z", "Estimate"]
  t_val = coef["X:Z", "t value"]
  #save to vectors
  ebeta_xz[i] = beta
  et_xz[i] = t_val
}
```

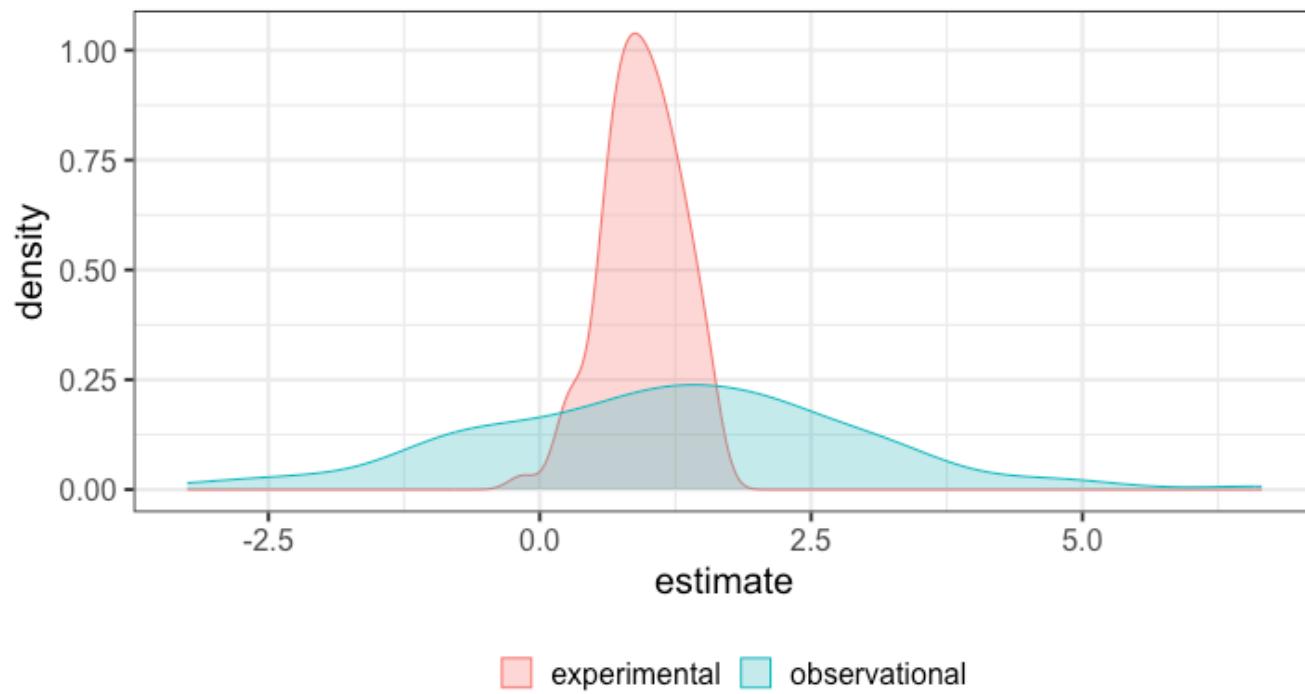
```
# for observational studies
obeta_xz = numeric(length = sim)
ot_xz = numeric(length = sim)
for(i in 1:sim){
  # simulate data
  X = rnorm(n = 100, mean=0,
  Z = rnorm(n = 100, mean=0,
  X = round(X/.2)*.2
  Z = round(Z/.2)*.2
  Y = 0 + 1*X + 1*Z + 1*X*Z
  rnorm(n = 100, mean = 0,
  #run model
  model = lm(Y ~ X*Z)
  coef = coef(summary(model))
  #extract coefficients
  beta = coef["X:Z", "Estimate"]
  t_val = coef["X:Z", "t value"]
  #save to vectors
  obeta_xz[i] = beta
  ot_xz[i] = t_val
}
```

```
mean(ebeta_xz)
```

```
## [1] 0.9440304
```

```
mean(obeta_xz)
```

```
## [1] 1.175444
```

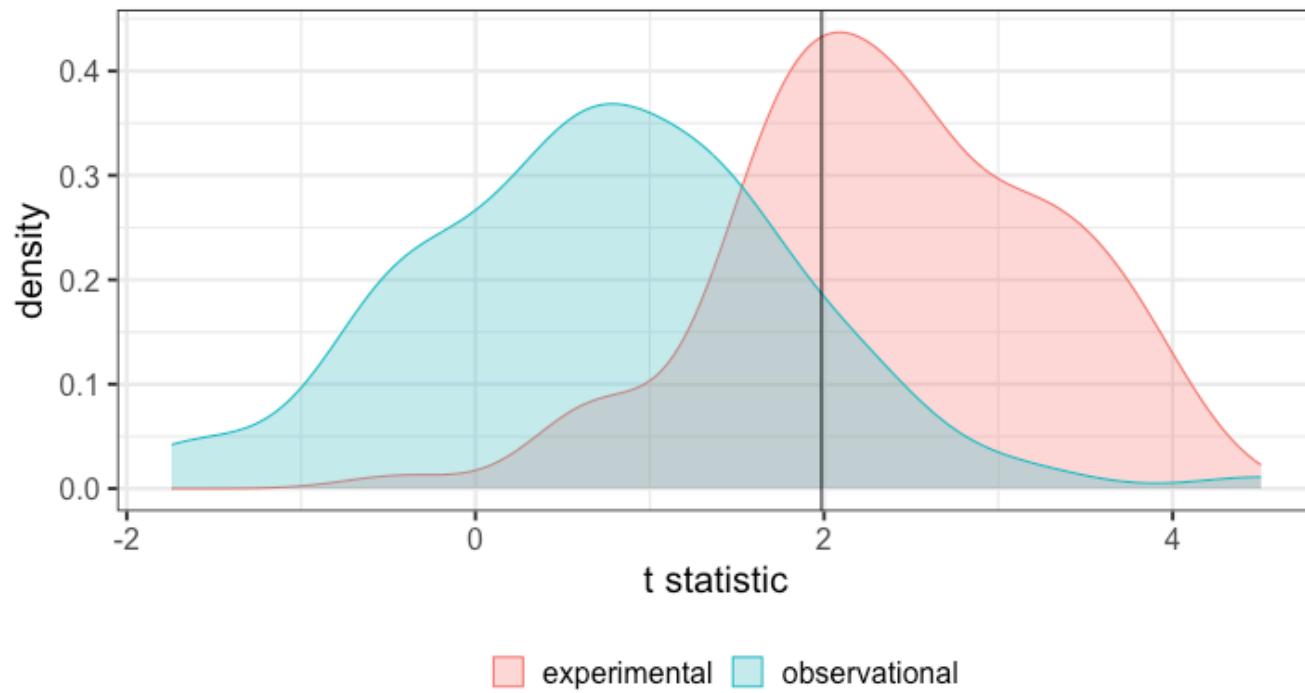


```
mean(et_xz)
```

```
## [1] 2.383435
```

```
mean(ot_xz)
```

```
## [1] 0.7411209
```



```
cv = qt(p = .975, df = 100-3-1)
esig = et_xz > cv
sum(esig)
```

```
## [1] 66
```

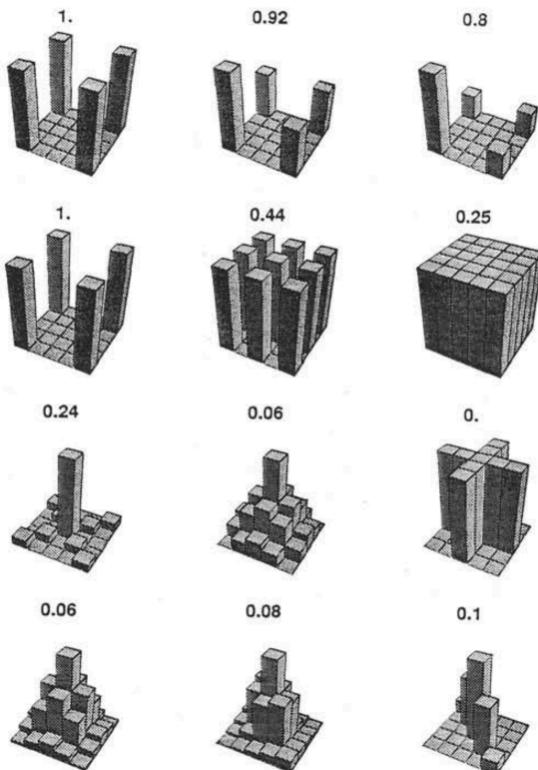
```
osig = ot_xz > cv
sum(osig)
```

```
## [1] 12
```

In our simulation, 66% of experimental studies were statistically significant, whereas only 12% of observational studies were significant.

Remember, we built our simulation based on data where there really is an interaction effect (i.e., the null is false).

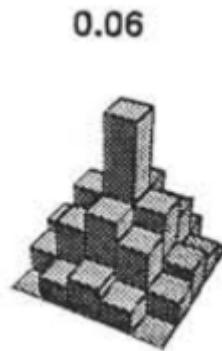
Efficiency



Efficiency = the ratio of the variance of XZ (controlling for X and Z) of a design to the best possible design (upper left corner). High efficiency is better; best efficiency is 1.

Efficiency

If the optimal design has N observations, then to have the same power, any other design needs to have $N * (1/\text{efficiency})$.



So a design with .06 efficiency needs $\frac{1}{.06} = 16.67$ times the sample size to detect the effect.

Efficiency

This particular point has been "rediscovered" as recently as 2018:

- **you need 16 times the sample size to detect an interaction as you need for a main effect of the same size.**
- This generalizes to higher-order interactions as well. If you have a three-way interaction, you need $16*16$ (256 times the number of people). More coming...

Observational studies: What NOT to do

Re-code X and Z into more extreme values
(e.g., median splits)

- while this increases variance in X and Z, it also increases measurement error

Collect a random sample and then only perform analyses on the sub sample with extreme values

- reduces sample size and also generalizability

Observational studies: What NOT to do

What can be done?

M&J suggest oversampling extremes and using weighted and unweighted samples

Experimental studies: What NOT to do

Be mean to field researchers

Forget about lack of external validity and generalizability

Ignore power when comparing interaction between covariate and experimental predictors (ANCOVA or multiple regression with categorical and continuous predictors)

Polynomials

Non-linear relationships

Linear lines often make bad predictions -- very few processes that we study actually have linear relationships. For example, effort had diminishing returns (e.g., log functions), or small advantages early in life can have significant effects on mid-life outcomes (e.g., exponential functions). In cases where the direction of the effect is constant but changing in magnitude, the best way to handle the data is to transform a variable (usually the outcome) and run linear analyses.

```
log_y = log(y)  
lm(log_y ~ x)
```

A small amount of anxiety is beneficial for performance on some tasks but too much is detrimental. When the **shape of the effect includes change(s) in direction**, then a **polynomial** term(s) may be more appropriate.

Polynomials are often a poor approximation for a non-linear effect. Correctly testing for non-linear effects usually requires (a) a lot of data and (b) making a number of assumptions about the data. Polynomial regression can be a useful tool for *exploratory* analysis and in cases when data are limited in terms of quantity and/or quality.

Polynomial regression

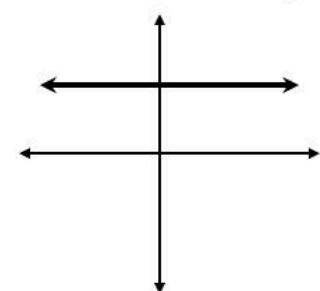
Polynomial regression is most often a form of hierarchical regression that systematically tests a series of higher order functions for a single variable.

Linear: $\hat{Y} = b_0 + b_1 X$

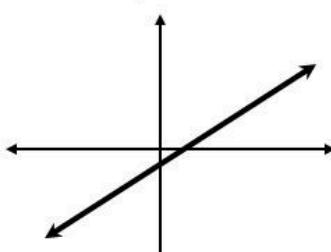
Quadratic: $\hat{Y} = b_0 + b_1 X + b_2 X^2$

Cubic: $\hat{Y} = b_0 + b_1 X + b_2 X^2 + b_3 X^3$

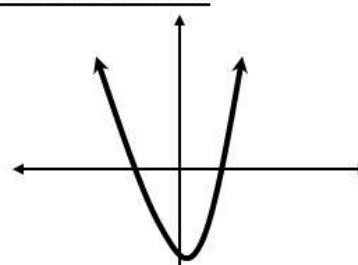
Graphs of Polynomial Functions:



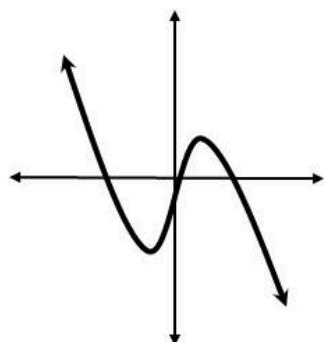
Constant Function
(degree = 0)



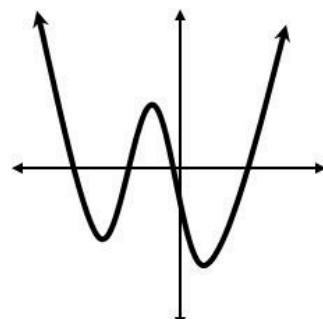
Linear Function
(degree = 1)



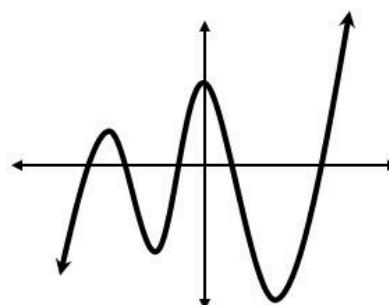
Quadratic Function
(degree = 2)



Cubic Function
(deg. = 3)

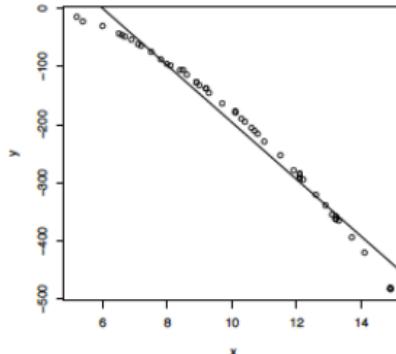


Quartic Function
(deg. = 4)

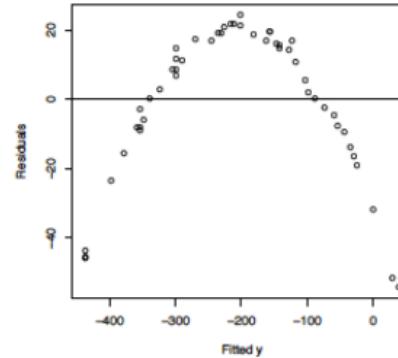


Quintic Function
(deg. = 5)

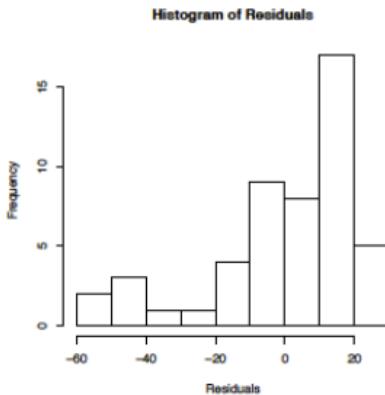
How do you know if something is up?



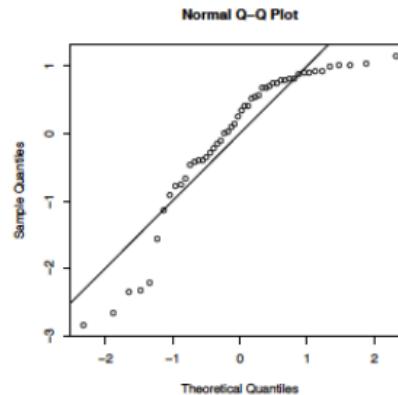
(a)



(b)



(c)



(d)

- (a) Scatterplot of the quadratic data with the OLS line. (b) Residual plot for the OLS fit.
(c) Histogram of the residuals. (d) NPP for the Studentized residuals.

<https://online.stat.psu.edu/stat501/book/export/html/962>

Can a team have too much talent?

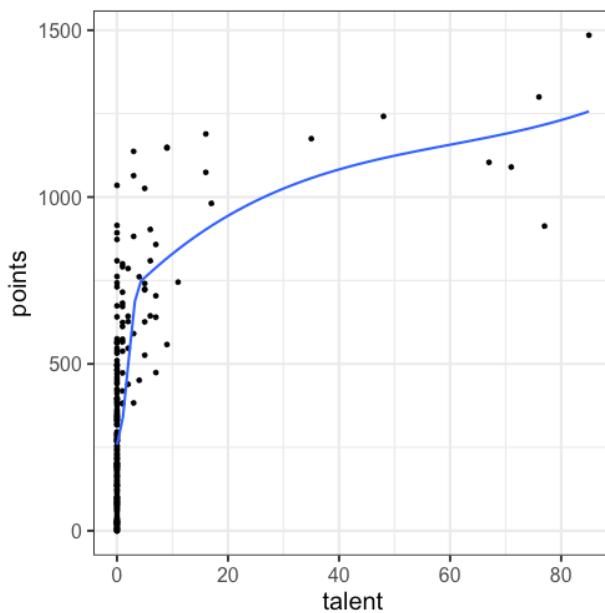
Do teams with too many talented players have poor intra-team coordination and perform worse than teams with a moderate amount of talent? They looked at 208 international football teams. Talent was the percentage of players during the 2010 and 2014 World Cup Qualifications phases who also had contracts with elite club teams. Performance was the number of points the team earned during these same qualification phases.

Swaab, R.I., Schaefer, M., Anicich, E.M., Ronay, R., and Galinsky, A.D. (2014). *The too-much-talent effect: Team interdependence determines when more talent is too much or not enough.* *Psychological Science* 25(8), 1581-1591.

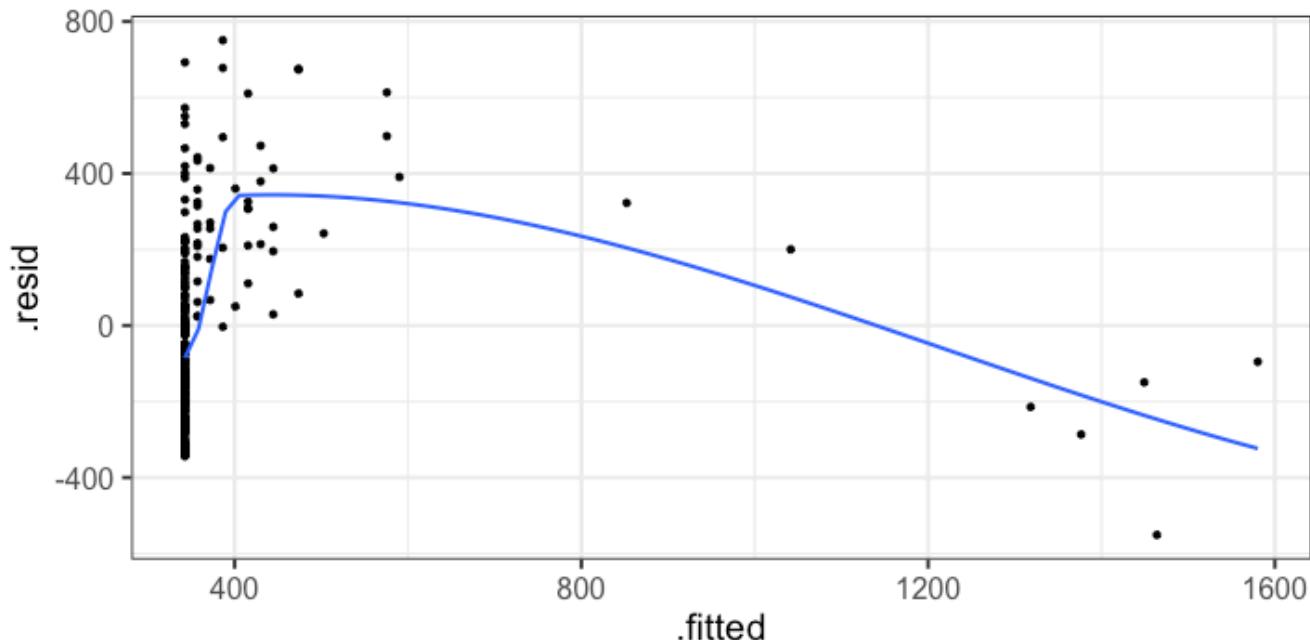
```
head(football)
```

```
##      country points talent
## 1       Spain    1485     85
## 2    Germany    1300     76
## 3     Brazil    1242     48
## 4   Portugal    1189     16
## 5 Argentina    1175     35
## 6 Switzerland   1149      9
```

```
ggplot(football, aes(x = talent,
geom_point() +
geom_smooth(se = F) +
theme_bw(base_size = 20))
```



```
mod1 = lm(points ~ talent, data = football)
library(broom)
aug1 = augment(mod1)
ggplot(aug1, aes(x = .fitted, y = .resid)) +
  geom_point() +
  geom_smooth(se = F) +
  theme_bw(base_size = 20)
```



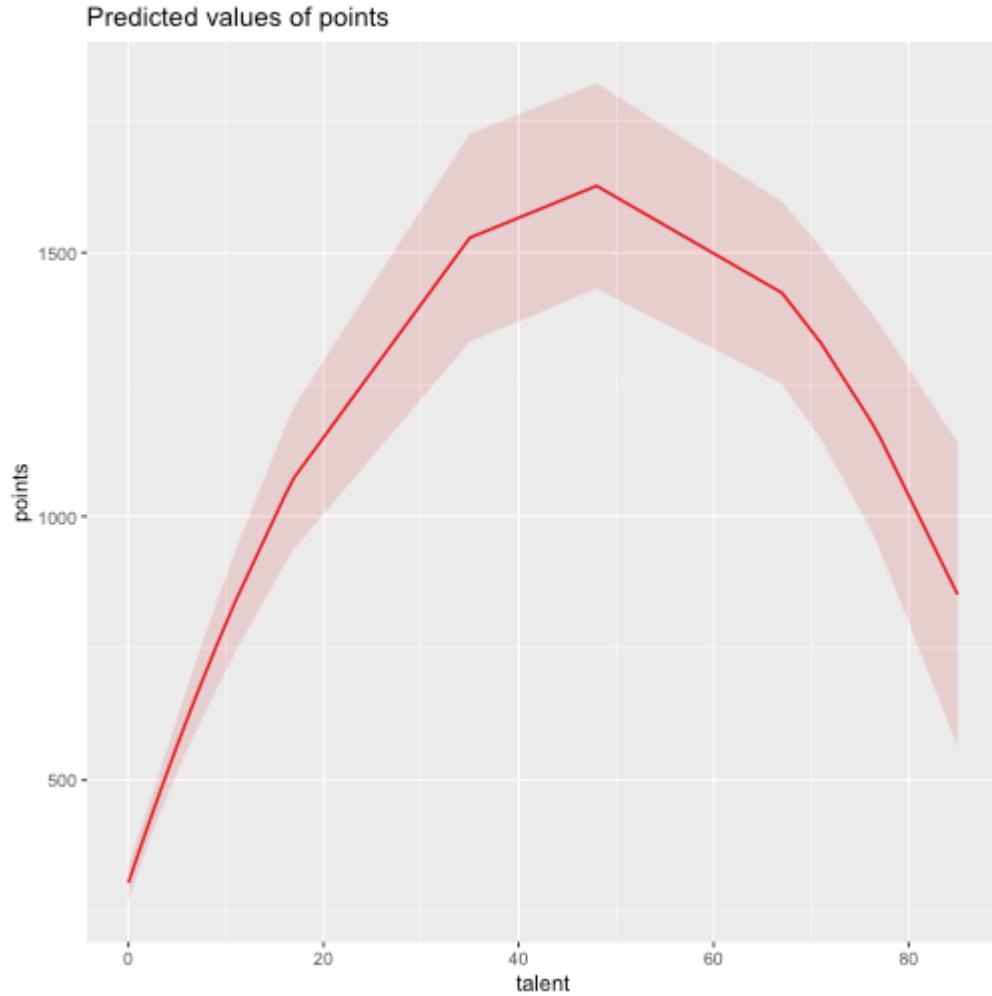
```
mod2 = lm(points ~ talent + I(talent^2), data = football)
```

The `I` stands for `identify`. It means that anything in `I()` should be treated according to its arithmetic meaning. Sometimes R gets confused when you include `^` and `~` in the same formula. This `I()` prevents such issues.

```
mod2 = lm(points ~ talent + I(talent^2), data = football)
summary(mod2)

##
## Call:
## lm(formula = points ~ talent + I(talent^2), data = football)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -384.66 -193.82  -35.34  152.11  729.66 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 305.34402   17.62668   17.323 < 2e-16 ***
## talent      54.89787    5.46864   10.039 < 2e-16 ***
## I(talent^2) -0.57022    0.07499   -7.604 1.01e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 236.3 on 205 degrees of freedom
## Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4592 
## F-statistic: 88.87 on 2 and 205 DF,  p-value: < 2.2e-16
```

```
library(sjPlot)
plot_model(mod2, type = "pred", terms = "talent")
```



Interpretation

The intercept is the predicted value of Y when $X = 0$

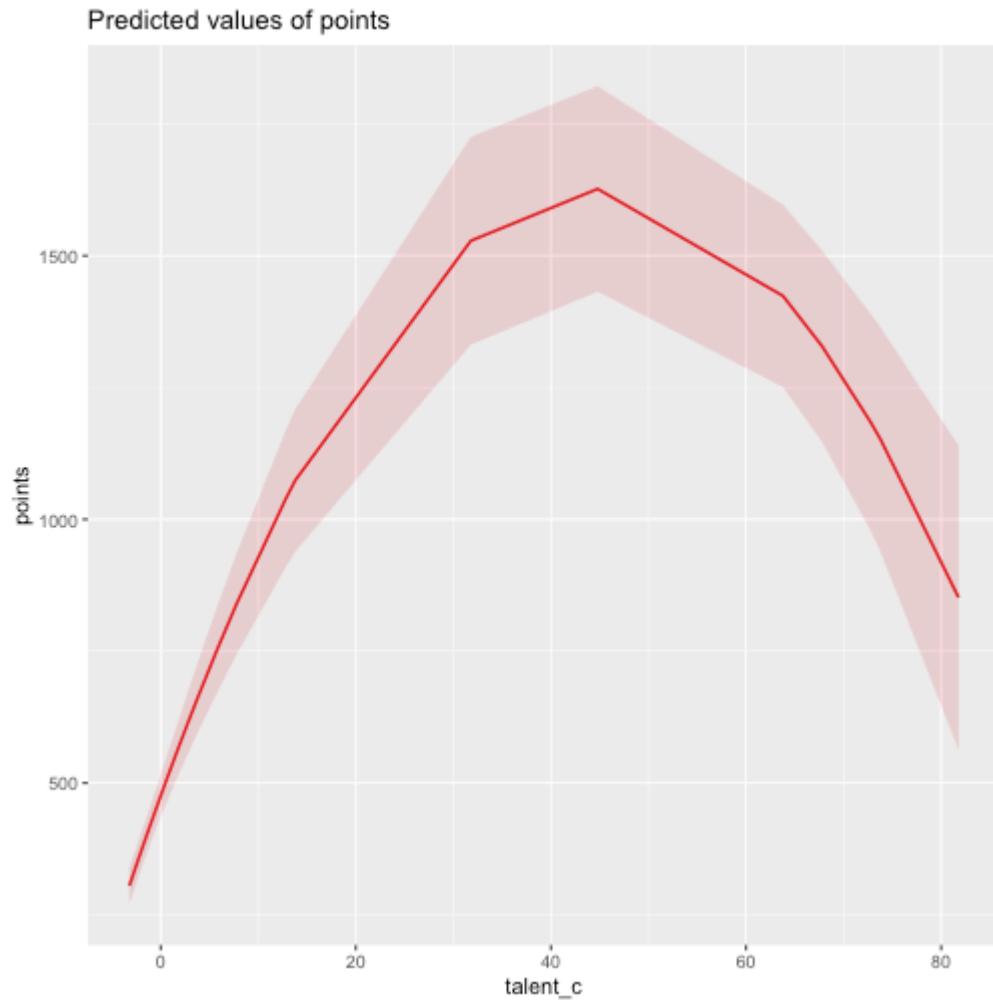
b_1 coefficient is the *tangent to the curve* when $X = 0$. In other words, this is the rate of change when X is equal to 0. If 0 is not a meaningful value on your X , you may want to center, as this will tell you the rate of change at the mean of X .

```
football$talent_c = football$talent - mean(football$talent)
mod2_c = lm(points ~ talent_c + I(talent_c^2), data = footbal
```

```
summary(mod2_c)
```

```
##  
## Call:  
## lm(formula = points ~ talent_c + I(talent_c^2), data = football)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max  
## -384.66 -193.82   -35.34  152.11  729.66  
##  
## Coefficients:  
##                 Estimate Std. Error t value Pr(>|t|)  
## (Intercept) 476.01572  19.94656 23.865 < 2e-16 ***  
## talent_c     51.22982   5.00212 10.242 < 2e-16 ***  
## I(talent_c^2) -0.57022   0.07499 -7.604 1.01e-12 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 236.3 on 205 degrees of freedom  
## Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4592  
## F-statistic: 88.87 on 2 and 205 DF,  p-value: < 2.2e-16
```

```
plot_model(mod2_c, type = "pred", terms = "talent_c")
```



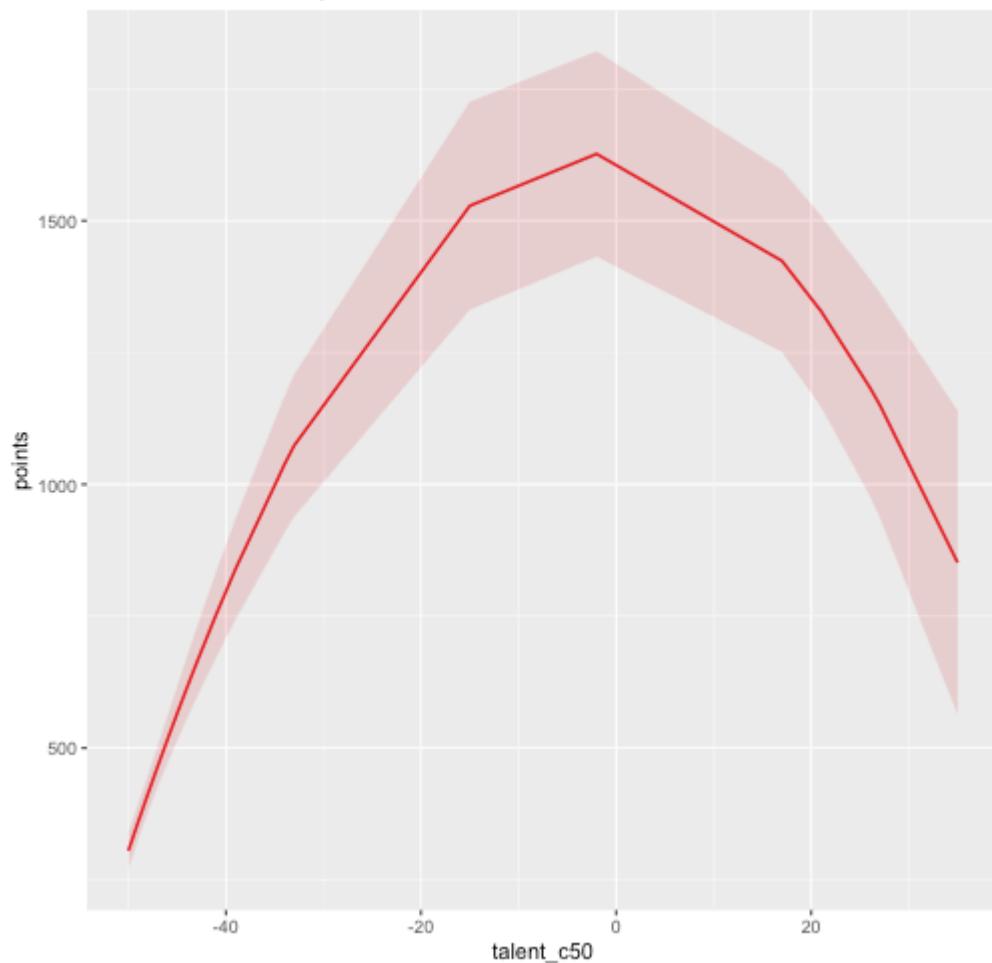
Or you can choose another value to center your predictor on, if there's a value that has a particular meaning or interpretation.

```
football$talent_c50 = football$talent - 50  
mod2_50 = lm(points ~ talent_c50 + I(talent_c50^2), data = fo
```

```
summary(mod2_50)
```

```
##  
## Call:  
## lm(formula = points ~ talent_c50 + I(talent_c50^2), data = football)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max  
## -384.66 -193.82   -35.34  152.11  729.66  
##  
## Coefficients:  
##                               Estimate Std. Error t value Pr(>|t|)  
## (Intercept)      1624.68872    97.18998  16.717 < 2e-16 ***  
## talent_c50        -2.12408     2.56568  -0.828    0.409  
## I(talent_c50^2)     -0.57022     0.07499  -7.604 1.01e-12 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 236.3 on 205 degrees of freedom  
## Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4592  
## F-statistic: 88.87 on 2 and 205 DF,  p-value: < 2.2e-16
```

Predicted values of points



Interpretation

The b_2 coefficient indexes the acceleration, which is how much the slope is going to change. More specifically, $2 \times b_2$ is the acceleration:

| the rate of change in b_1 for a 1-unit change in X

You can use this to calculate the slope of the tangent line at any value of X you're interested in:

$$b_1 + (2 \times b_2 \times X)$$

```
tidy(mod2)
```

```
## # A tibble: 3 × 5
##   term      estimate std.error statistic p.value
##   <chr>      <dbl>     <dbl>     <dbl>     <dbl>
## 1 (Intercept) 305.       17.6      17.3  5.22e-42
## 2 talent      54.9       5.47      10.0  1.53e-19
## 3 I(talent^2) -0.570     0.0750    -7.60 1.01e-12
```

At X = 10

```
54.9 + (2 * -.570 * 10)
```

```
## [1] 43.5
```

At X = 70

```
54.9 + (2 * -.570 * 70)
```

```
## [1] -24.9
```

Polynomials are interactions

A term for X^2 is a term for $X \times X$ or the multiplication of two independent variables holding the same values.

```
football$talent_2 = football$talent*football$talent  
tidy(lm(points ~ talent + talent_2, data = football))
```

```
## # A tibble: 3 × 5  
##   term      estimate std.error statistic p.value  
##   <chr>      <dbl>     <dbl>     <dbl>    <dbl>  
## 1 (Intercept) 305.       17.6      17.3  5.22e-42  
## 2 talent      54.9       5.47     10.0  1.53e-19  
## 3 talent_2    -0.570     0.0750    -7.60 1.01e-12
```

Polynomials are interactions

Put another way:

$$\hat{Y} = b_0 + b_1 X + b_2 X^2$$

$$\hat{Y} = b_0 + \frac{b_1}{2} X + \frac{b_1}{2} X + b_2(X \times X)$$

The interaction term in another model would be interpreted as "how does the slope of X change as I move up in Z?" -- here, we ask "how does the slope of X change as we move up in X?"

When should you use polynomial terms?

You may choose to fit a polynomial term after looking at a scatterplot of the data or looking at residual plots. A U-shaped curve may be indicative that you need to fit a quadratic form -- although, as we discussed before, you may actually be measuring a different kind of non-linear relationship.

Polynomial terms should mostly be dictated by theory -- if you don't have a good reason for thinking there will be a change in sign, then a polynomial is not right for you.

Three-way interactions and beyond

Three-way interactions (regression)

Regression equation

$$\hat{Y} = b_0 + b_1X + b_2Z + b_3W + b_4XZ + b_5XW + b_6ZW + b_7XZW$$

The three-way interaction qualifies the three main effects (and any two-way interactions).

Like a two-way interaction, the three-way interaction is a conditional effect. And it is symmetrical, meaning there are several equally correct ways of interpreting it.

How do we describe a 3-way ANOVA?

A two-way ($A \times B$) interaction means that the magnitude of one main effect (e.g., A main effect) depends on levels of the other variable (B). But, it is equally correct to say that the magnitude of the B main effect depends on levels of A. In regression, we refer to these as **conditional effects** and in ANOVA, they are called **simple main effects**.

A three-way interaction means that the *magnitude of one two-way interaction (e.g., $A \times B$) depends on the levels of the remaining variable (C)*.

A three-way interaction means that the *magnitude of one two-way interaction (e.g., A x B) depends on the levels of the remaining variable (C).*

It is equally correct to say that the magnitude of the A x C interaction depend on levels of B. Or, that the magnitude of the B x C interaction depends on levels of A. These are known as **simple interaction effects**.

```
psych::describe(stress_data, fast = T)
```

	vars	n	mean	sd	median	min	max	range	skew	kurtosis	se
## gender	1	150	NaN	NA	NA	Inf	-Inf	-Inf	NA	NA	NA
## bad_day	2	150	2.95	1.29	3	1	5	4	0.03	-1.11	0.11
## talk	3	150	2.57	1.20	2	1	5	4	0.32	-0.94	0.10
## stress	4	150	30.15	10.00	30	1	51	50	-0.03	-0.56	0.82

```
table(stress_data$gender)
```

```
##  
## female male  
## 67 83
```

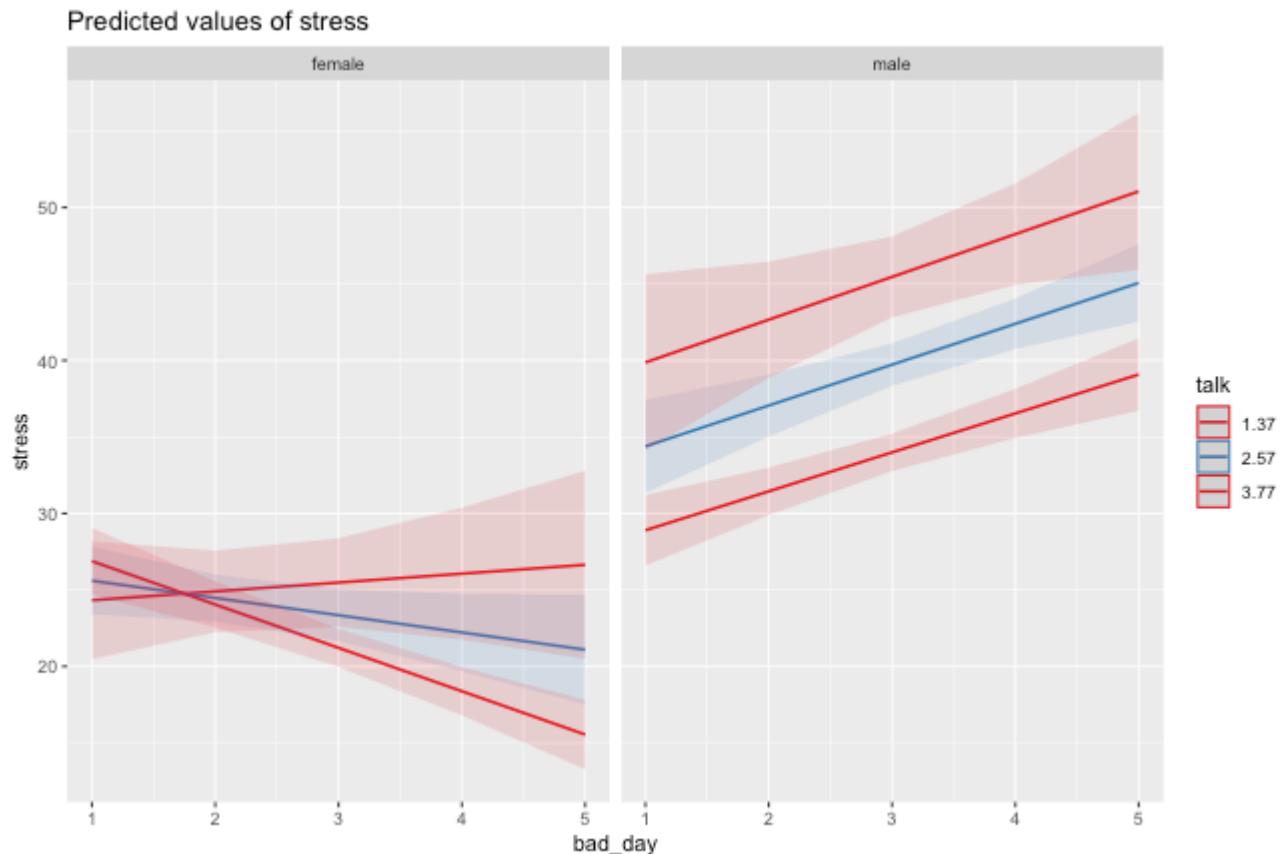
```
mod_stress = lm(stress ~ bad_day*talk*gender, data = stress_d
summary(mod_stress)

## 
## Call:
## lm(formula = stress ~ bad_day * talk * gender, data = stress_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.6126  -3.2974   0.0671   3.1129  10.7774
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)              20.3385    4.5181   4.502 1.39e-05 ***
## bad_day                  2.5273    1.6596   1.523  0.13003    
## talk                     2.4870    1.3234   1.879  0.06227    
## gendermale               -0.1035    5.7548  -0.018  0.98568    
## bad_day:talk              -1.4220    0.4564  -3.116  0.00222 **  
## bad_day:gendermale        -0.1244    2.0069  -0.062  0.95067    
## talk:gendermale            1.9797    2.2823   0.867  0.38718    
## bad_day:talk:gendermale   1.5260    0.7336   2.080  0.03931 *  
## ---                        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.733 on 142 degrees of freedom
```

```
library(reghelper)
simple_slopes(mod_stress)
```

##	bad_day	talk	gender	Test	Estimate	Std. Error	t value	df	Pr(> t)
## 1	1.661455	1.365903	sstest		5.8571	1.7437	3.3590	142	0.001
## 2	2.953333	1.365903	sstest		8.3892	1.5671	5.3535	142	3.386
## 3	4.245211	1.365903	sstest		10.9214	2.5606	4.2651	142	3.627
## 4	1.661455	2.566667	sstest		11.2788	1.4571	7.7406	142	1.687
## 5	2.953333	2.566667	sstest		16.1781	1.0985	14.7274	142	< 2.2
## 6	4.245211	2.566667	sstest		21.0775	1.6775	12.5647	142	< 2.2
## 7	1.661455	3.767431	sstest		16.7004	2.3888	6.9912	142	9.806
## 8	2.953333	3.767431	sstest		23.9670	1.4830	16.1608	142	< 2.2
## 9	4.245211	3.767431	sstest		31.2335	2.0578	15.1784	142	< 2.2
## 10	1.661455	sstest	female		0.1245	0.7221	0.1724	142	0.863
## 11	2.953333	sstest	female		-1.7125	0.5996	-2.8558	142	0.004
## 12	4.245211	sstest	female		-3.5495	0.9450	-3.7561	142	0.000
## 13	sstest	1.365903	female		0.5850	1.0725	0.5455	142	0.586
## 14	sstest	2.566667	female		-1.1224	0.6181	-1.8160	142	0.071
## 15	sstest	3.767431	female		-2.8299	0.4630	-6.1115	142	9.006
## 16	1.661455	sstest	male		4.6396	1.0195	4.5510	142	1.137
## 17	2.953333	sstest	male		4.7741	0.6463	7.3862	142	1.177
## 18	4.245211	sstest	male		4.9085	0.9474	5.1813	142	7.418
## 19	sstest	1.365903	male		2.5451	0.5036	5.0533	142	1.315
## 20	sstest	2.566667	male		2.6700	0.6116	4.3655	142	2.427
## 21	sstest	3.767431	male		2.7950	1.2025	2.3244	142	0.021

```
figures = plot_model(mod_stress, type = "int", mdrt.values =  
figures[[4]]
```



As a reminder, centering will change all but the highest-order terms in a model.

```
stress_data = stress_data %>%
  mutate(bad_day_c = bad_day - mean(bad_day),
        talk_c = talk - mean(talk))
newmod = lm(stress ~ bad_day_c*talk_c*gender, data = stress_d
```

```
tidy(mod_stress)
```

```
## # A tibble: 8 × 5
##   term                estimate std.error statistic p.value
##   <chr>              <dbl>     <dbl>      <dbl>    <dbl>
## 1 (Intercept)        20.3       4.52       4.50    0.0000139
## 2 bad_day            2.53       1.66       1.52    0.130
## 3 talk               2.49       1.32       1.88    0.0623
## 4 gendermale         -0.104      5.75      -0.0180  0.986
## 5 bad_day:talk       -1.42       0.456      -3.12   0.00222
## 6 bad_day:gendermale -0.124      2.01      -0.0620  0.951
## 7 talk:gendermale    1.98       2.28       0.867   0.387
## 8 bad_day:talk:gendermale 1.53       0.734      2.08    0.0393
```

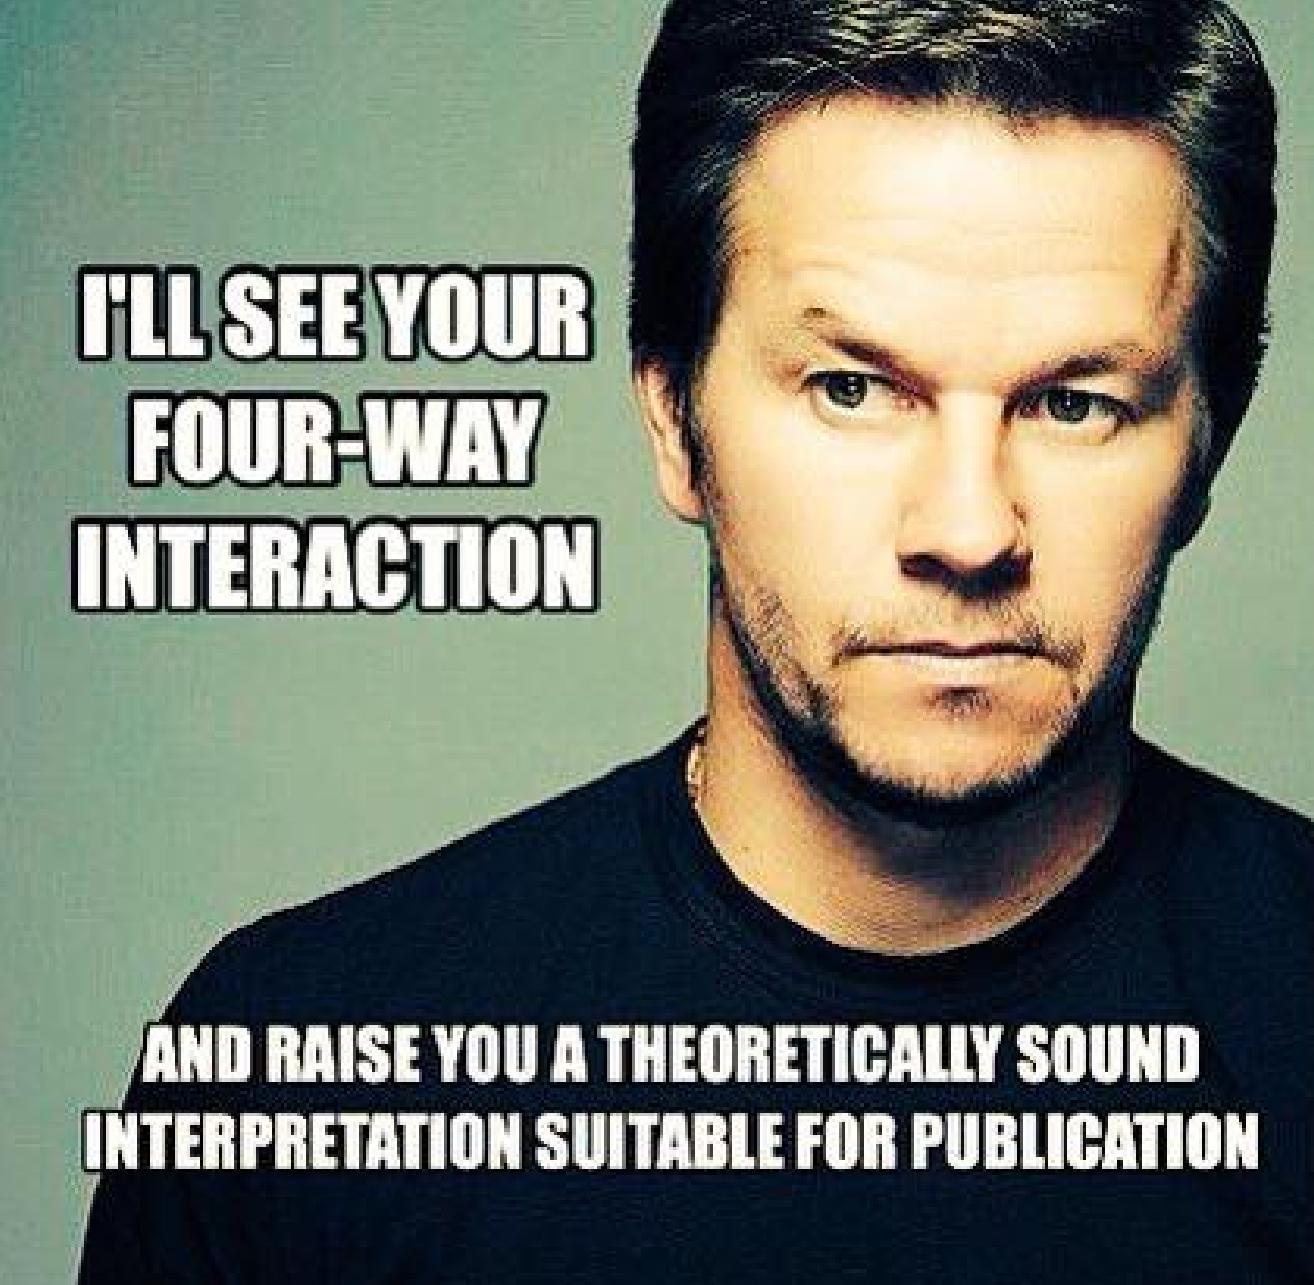
```
tidy(newmod)
```

## # A tibble: 8 × 5	## term	## <chr>	estimate	std.error	statistic	p.value
## 1 (Intercept)			23.4	0.845	27.7	3.89e-59
## 2 bad_day_c			-1.12	0.618	-1.82	7.15e- 2
## 3 talk_c			-1.71	0.600	-2.86	4.94e- 3
## 4 gendermale			16.2	1.10	14.7	2.20e-30
## 5 bad_day_c:talk_c			-1.42	0.456	-3.12	2.22e- 3
## 6 bad_day_c:gendermale			3.79	0.870	4.36	2.47e- 5
## 7 talk_c:gendermale			6.49	0.882	7.36	1.38e-11
## 8 bad_day_c:talk_c:gendermale			1.53	0.734	2.08	3.93e- 2

Four-way?

$$\begin{aligned}\hat{Y} = & b_0 + b_1 X + b_2 Z + b_3 W + b_4 Q + b_5 XW \\& + b_6 ZW + b_7 XZ + b_8 QX + b_9 QZ + b_{10} QW \\& + b_{11} XZQ + b_{12} XZW + b_{13} XWQ + b_{14} ZWQ \\& + b_{15} XZWQ\end{aligned}$$

3-way (and higher) interactions are incredibly difficult to interpret, in part because they represent incredibly complicated processes. If you have a solid theoretical rationale for conducting a 3-way interaction, be sure you've collected enough subjects to power your test.



**I'LL SEE YOUR
FOUR-WAY
INTERACTION**

**AND RAISE YOU A THEORETICALLY SOUND
INTERPRETATION SUITABLE FOR PUBLICATION**

Especially with small samples, three-way interactions may be the result of a few outliers skewing a regression line. If you have stumbled upon a three-way interaction during exploratory analyses, **be careful**. This is far more likely to be a result of over-fitting than uncovering a true underlying process.

Use at least one nominal moderator (ideally with only 2 levels), instead of all continuous moderators. This allows you to examine the 2-way interaction at each level of the nominal moderator, esp if one moderator is experimenter manipulated, which increases the likelihood of balanced conditions.

Next time...

Logistic Regression

Catch Up/Flex -- if nothing, then Ke will run the review session in class

March 27th = Exam 2!!!!!!!

April 1st = discussion of Yarkoni & Westfall paper. You 100% need to read it. Not optional.

(everything after this slide is extra or code-based; feel free to play around, but you won't be tested on it)

```
library(car)
fit = lm(Time~Speed*Noise, data = Data)
summary(aov(fit))
```

```
##           Df  Sum Sq Mean Sq F value    Pr(>F)
## Speed       2 2805871 1402936 109.397 < 2e-16 ***
## Noise       2   341315   170658  13.307 4.25e-06 ***
## Speed:Noise 4   295720    73930   5.765 0.000224 ***
## Residuals 171 2192939    12824
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Anova(fit, type = 2)
```

```
## Anova Table (Type II tests)
##
## Response: Time
##           Sum Sq Df  F value    Pr(>F)
## Speed       2805871  2 109.3975 < 2.2e-16 ***
## Noise       341315  2  13.3075 4.252e-06 ***
## Speed:Noise 295720  4   5.7649 0.0002241 ***
## Residuals 2192939 171
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
fit = lm(Time~Speed*Noise, data = Data)
summary(fit)
```

```
## 
## Call:
## lm(formula = Time ~ Speed * Noise, data = Data)
## 
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -316.23  -70.82     4.99    79.87   244.40 
## 
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)                 630.72    25.32  24.908 < 2e-16 ***
## SpeedMedium                -105.44    35.81  -2.944  0.00369 **  
## SpeedFast                  -301.45    35.81  -8.418  1.49e-14 ***
## NoiseControllable          -54.05    35.81  -1.509  0.13305    
## NoiseUncontrollable        -36.28    35.81  -1.013  0.31243    
## SpeedMedium:NoiseControllable  21.48    50.64   0.424  0.67201    
## SpeedFast:NoiseControllable  12.01    50.64   0.237  0.81287    
## SpeedMedium:NoiseUncontrollable -184.39   50.64  -3.641  0.00036 *** 
## SpeedFast:NoiseUncontrollable -24.84    50.64  -0.490  0.62448    
## ---                        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
## 
## Residual standard error: 113.2 on 171 degrees of freedom
```

```

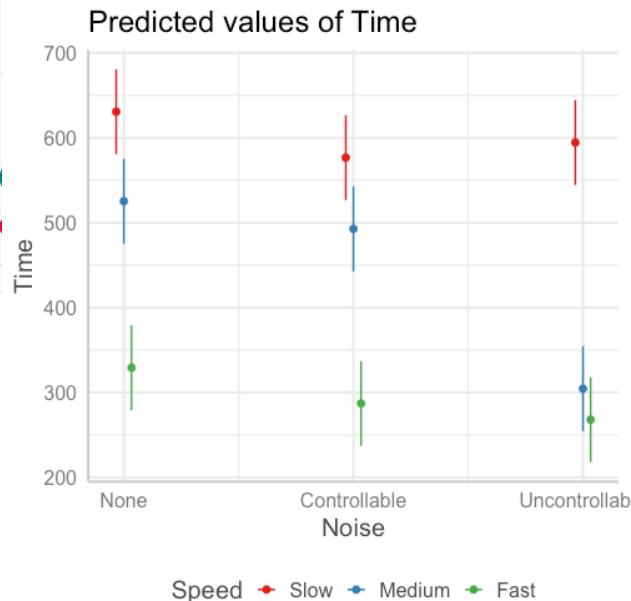
library(emmeans)
fit.grid = ref_grid(fit)
pairs(fit.grid, adjust = "holm")

```

	estimate	SE	df	t.ratio
## contrast				
## Slow None - Medium None	105.4	35.8	171	2.944
## Slow None - Fast None	301.4	35.8	171	8.418
## Slow None - Slow Controllable	54.1	35.8	171	1.509
## Slow None - Medium Controllable	138.0	35.8	171	3.854
## Slow None - Fast Controllable	343.5	35.8	171	9.592
## Slow None - Slow Uncontrollable	36.3	35.8	171	1.013
## Slow None - Medium Uncontrollable	326.1	35.8	171	9.106
## Slow None - Fast Uncontrollable	362.6	35.8	171	10.124
## Medium None - Fast None	196.0	35.8	171	5.473
## Medium None - Slow Controllable	-51.4	35.8	171	-1.435
## Medium None - Medium Controllable	32.6	35.8	171	0.910
## Medium None - Fast Controllable	238.1	35.8	171	6.648
## Medium None - Slow Uncontrollable	-69.2	35.8	171	-1.931
## Medium None - Medium Uncontrollable	220.7	35.8	171	6.162
## Medium None - Fast Uncontrollable	257.1	35.8	171	7.180
## Fast None - Slow Controllable	-247.4	35.8	171	-6.908
## Fast None - Medium Controllable	-163.4	35.8	171	-4.564
## Fast None - Fast Controllable	42.0	35.8	171	1.174
## Fast None - Slow Uncontrollable	-265.2	35.8	171	-7.405
## Fast None - Medium Uncontrollable	24.7	35.8	171	0.689
## Fast None - Fast Uncontrollable	61.1	35.8	171	1.707

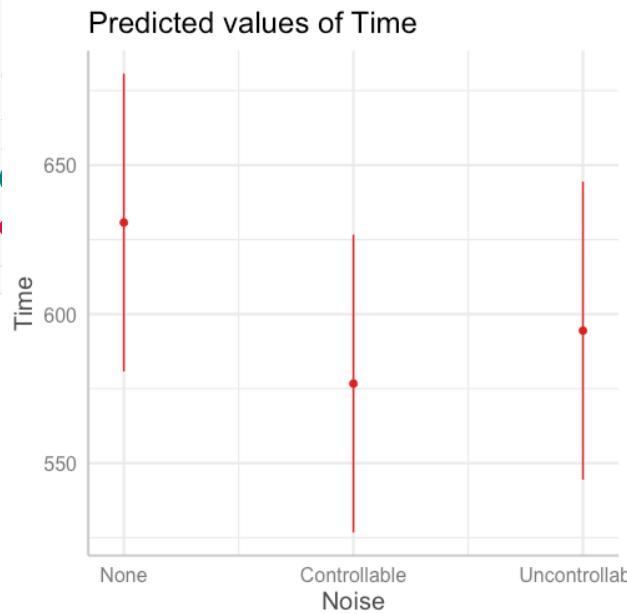
In `sjPlot()`, predicted values are the expected value of Y given all possible values of X, **at specific values of M**. If you don't give it all of M, it will choose every possible value.

```
library(sjPlot)
plot_model(fit, type = "pred"
           terms = c("Noise"
                     "Speed"))
theme_sjplot(base_size = 21)
theme(legend.position = "bottom")
```



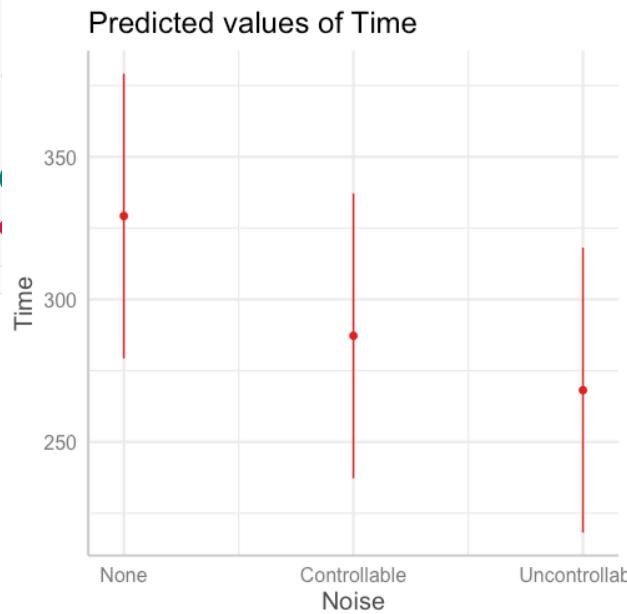
In `sjPlot()`, predicted values are the expected value of Y given all possible values of X, **at specific values of M**. If you don't specify levels of M, it will choose the lowest possible value.

```
library(sjPlot)
plot_model(fit, type = "pred"
           terms = c("Noise"))
theme_sjplot(base_size = 21)
theme(legend.position = "bottom")
```



In `sjPlot()`, predicted values are the expected value of Y given all possible values of X, **at specific values of M**. If you don't specify levels of M, it will choose the lowest possible value.

```
library(sjPlot)
plot_model(fit, type = "pred"
           terms = c("Noise")
           theme_sjplot(base_size = 21)
           theme(legend.position = "bottom"))
```



In `sjPlot()`, estimated marginal means are the expected value of Y given all possible values of X, **ignoring M**.

```
library(sjPlot)
plot_model(fit, type = "emm"
           terms = c("Noise"
                     theme_sjplot(base_size = 20)
                     theme(legend.position = "bottom"))
```

